

▸ Tp5 Dimitris

Non-blind watermark detection

Exercise 1

Read the image cameraman.tif in. It will serve as host image x . For given hypothesis:

$$H0: v = x + z$$

$$H1: v = x + w + z$$

where x is the host image, v is the marked image, w is the watermark and z is additive white Gaussian noise, i.e. $Z \sim N(0, \sigma_{noise}^2 I)$

```
1 import matplotlib.image as mpimg; import matplotlib.pyplot as plt; import numpy
2
3 url = "https://github.com/partizanos/multimedia_security/raw/master/multimedia_
4 response = requests.get(url)
5 im = Image.open(BytesIO(response.content))
6 X = np.array(im)
7 N1, N2 = X.shape
8 N = int(N1*N2)
```

Hypothesis and watermark definition as on TP4

```
1 def imageShape(img):
2     N1, N2 = X.shape
3     assert X.shape == (256, 256)
4     return int(N1*N2)
5
6
7 def watermark(gamma, img, theta):
8     assert type(gamma) == int
9     N1, N2 = img.shape
10    size = imageShape(img)
11    assert type(size) == int
12    indices=[]
13
14    # Generate a matrix w' of size x with uniform distributed values {-1, 1}.
15    # The magnitude of these two values governs the watermark strength.
16    w_prime = np.random.uniform(-gamma, gamma, size)
17    w_prime = w_prime.reshape((N1, N2))
18
19    # Randomly sample from matrix w' with a given density  $\theta N = 0.5$ 
20    for i in range(int(size * theta)):
21        indices.append((np.random.randint(N1), np.random.randint(N2)))
22
23    w = np.zeros((N1, N2))
24    for tup in indices:
25        w[tup] = w_prime[tup]
26    return w
27
28
29
30 def get_Hypothesis(X, sigma_noise, gamma, theta):
31     #  $V = X + Z$ 
32     N1, N2 = X.shape
33     Z = np.random.normal(0, sigma_noise, imageShape(X)).reshape((N1, N2))
34     V_H0 = X + Z
35     w = watermark(gamma, X, theta)
```

```

36 | V_H1 = V_H0 + w
37 | return V_H0, V_H1, w
38 |
39 | def rho_non_blind(v, w, X, N ):
40 |     w_non_blind = v - X
41 |     return np.sum(w.reshape((N))*w_non_blind.reshape((N))) / (N)
42 |

```

For hypothesis H_1 and ρ^{H_1} the watermark w is generated J times with a fixed strength $\gamma = \pm 1.2$ and a fixed density $\theta_N = 0.1$. The noise realization z is again fixed with $\sigma_{noise} = 50$.

Fill out Table 1 with all results. Note that obviously only the noise and not the watermark has influence on hypothesis H_0 , so the relevant cells have been grayed out.

Determine rho mu and sigma parameter for H_0 and H_1

```

1 | # sigmas =[50, 100];
2 | # densities = [0.1, 0.3];
3 | # gammas = [1, 5];
4 |
5 | # sigma_noise = 50
6 | # theta = 0.1
7 | # gamma = 1
8 |
9 |
10 | def experiment(X, sigma_noise, gamma, theta, J ):
11 |     rho_h0_list = []
12 |     rho_h1_list = []
13 |
14 |     for j in range(J):
15 |         vh0, vh1, w = get_Hypothesis(X, sigma_noise, gamma, theta)
16 |         rho_h0_list.append(rho_non_blind(vh0, w, X, N ))
17 |         rho_h1_list.append(rho_non_blind(vh1, w, X, N ))
18 |
19 |     rho_h0_list = np.array(rho_h0_list)
20 |     rho_h1_list = np.array(rho_h1_list)
21 |
22 |     mu_h0 = np.sum(rho_h0_list)/J
23 |     mu_h1 = np.sum(rho_h1_list)/J
24 |
25 |     sigma_h0 = (np.sum(rho_h0_list - mu_h0)**2) / J
26 |     sigma_h1 = (np.sum(rho_h1_list - mu_h1)**2) / J
27 |
28 |     print("sigma noise: ", sigma_noise)
29 |     print("theta: ", theta)
30 |     print("gamma: ", gamma)
31 |     print("### mu|h0: ", mu_h0, ", : mu|h1 ", mu_h1, " sigma_h0: ", sigma_h0, "sig
32 |     return mu_h0, mu_h1, sigma_h0, sigma_h1

```



```

1 | sigma_noise, gamma, theta = 50, 5, 0.3
2 | vh0, vh1, w = get_Hypothesis(X, sigma_noise, gamma, theta)
3 |
4 | fig, ax = plt.subplots(
5 |     ncols=2,
6 |     sharex=True,
7 |     sharey=True,
8 |     figsize=(18, 18)
9 | )
10 |
11 | ax[0].imshow(vh0, 'gray')
12 | ax[0].set_title("h0")
13 | ax[1].imshow(vh1, 'gray')
14 | ax[1].set_title("h1")

```



Calculate and display the Receiver Operating Characteristic (ROC) curve for the binary threshold test following the above mentioned experiment set up. The detection threshold is denoted with $T_{\rho_{non-blind}}$.

```

1 from scipy.stats import norm
2 sigma_noise, gamma, theta = 100, 5, 0.3
3
4 [mu_H0, mu_H1, var_H0, var_H1] = experiment(X, sigma_noise, gamma, theta, J );
5 threshold = np.linspace(mu_H0-var_H0,mu_H1+var_H1,12);
6 PFA=[]
7 PMISS=[]
8 for i in range(len(threshold)):
9     PFA.append(1 - norm.cdf(threshold[i], mu_H0, var_H0))
10    PMISS.append(norm.cdf(threshold[i], mu_H1, var_H1))
11
12 plt.plot(PFA); plt.show(); plt.title("PFA")
13 plt.plot(PMISS); plt.show(); plt.title("PMISS")
14 plt.plot(PFA,PMISS); plt.show(); plt.title("ROC")

```



```
1 |  
2 | sigmas =[50, 100];  
3 | densities = [0.1, 0.3];
```

```
4 gammas = [1, 5];
5 J = 100
6 for sigma in sigmas:
7     for theta in densities:
8         for gamma in gammas:
9             print("Running experiment with parameters: ", sigma, theta, gamma)
10             mu_h0, mu_h1, sigma_h0, sigma_h1 = experiment(X, sigma, gamma, theta, J)
11
```



What can you conclude about non-blind watermark detection given the strength of the watermark and the noise variance?

Answer

- For higher values of theta watermark mu is significantly higher.
- For sigma 100 the standard deviation was 2 orders of magnitude higher than for sigma 50.

Double-click (or enter) to edit

